

CDCI 567-Machine Learning Assignment-1

1 Density Estimation

1.1 MLE

1.1.1 Beta($\alpha, 1$)

x_1, x_2, \dots, x_n are samples of $Beta(\alpha, 1)$ distribution.

In this case, pdf of $x(i), i = 1, \dots, n$ is

$$f(x_i|\alpha, 1) = \frac{x^{\alpha-1}}{\beta(\alpha, 1)} = \frac{x^{\alpha-1}}{\int_0^1 t^{\alpha-1} dt} = \alpha x^{\alpha-1}$$

Likelihood function and log likelihood function:

$$L(\alpha, 1|x) = \prod_{i=1}^n \alpha x^{\alpha-1}$$

$$l(\alpha, 1|x) = \log(\alpha^n) + \log\left(\prod_{i=1}^n x_i^{\alpha-1}\right) = n\log(\alpha) + \sum_{i=1}^n (\alpha - 1)\log(x_i)$$

Derivative of $l(\alpha, 1|x)$:

$$\frac{\partial l(\alpha, 1|x)}{\partial \alpha} = \frac{n}{\alpha} + \sum_{i=1}^n \log(x_i) = 0$$

$$\frac{\partial^2 l(\alpha, 1|x)}{\partial \alpha^2} = -\frac{n}{\alpha^2} < 0$$

$$\hat{\alpha} = -\frac{n}{\sum_{i=1}^n \log(x_i)}$$

1.1.2 Normal($\theta, \text{diag}(\theta)$)

x_1, x_2, \dots, x_n are samples of $Normal(\theta, \text{diag}(\theta))$ distribution. In this case, pdf of $x_i, i = 1, \dots, n$ is

$$f(x_i|\theta, \text{diag}(\theta)) = \frac{1}{\sqrt{2\pi \text{diag}(\theta)}} \exp\left(-\frac{(x_i - \theta)}{2\theta}\right)$$

Likelihood function and log likelihood function:

$$L(\theta, \text{diag}(\theta)|x_i) = (2\pi)^{-\frac{d}{2}} |\text{diag}(\theta)|^{-\frac{1}{2}} \exp\left(-\frac{1}{2}(x - \theta)^T \text{diag}(\theta)^{-1}(x - \theta)\right)$$

$$|diag(\theta)| = \theta^d$$

$$l(\theta, diag(\theta)|x_i) = -\frac{d}{2}\log(2\pi\theta) - \sum_{i=1}^n \frac{1}{2} \frac{(x_i - \theta_i)^T (x_i - \theta_i)}{\theta_i}$$

Derivative of $l(\theta, diag(\theta)|x_i)$:

$$\frac{\partial l(\theta, diag(\theta)|x_i)}{\partial \theta} = -\frac{dn}{2\theta} - \frac{n}{2} + \frac{1}{2\theta^2} \sum_{i=1}^n x_i^2 = 0$$

$$\hat{\theta} = -\frac{d}{2} + \frac{1}{2} \sqrt{d + \frac{4}{n} \sum_{i=1}^n x_i^2}$$

1.2 MLE in Linear Regression

1.2.1 w_0 and w

According to lecture pdf:

$$\log P(D) = -\frac{1}{2} \left(\frac{1}{\sigma^2} \sum_n [y_n - (w_0 + w^T x_n)]^2 \right)$$

To maximize the log likelihood, we need to minimize:

$$\sum_n [y_n - (w_0 + w^T x_n)]^2 = \sum_n [y_n^2 - 2y_n w_0 - 2y_n w^T x_n + w_0^2 + 2w_0 w^T x_n + (w^T x_n)^2]$$

$$\frac{\partial \sum_n [y_n - (w_0 + w^T x_n)]^2}{\partial w_0} = -2 \sum y_n + 2Nw_0 + 2 \sum w^T x_n = 0$$

$$\hat{w}_0 = \frac{\sum y_n - \sum w^T x_n}{N} = \bar{y} - w^T \bar{x}$$

$RSS(w)$ in matrix form:

$$RSS(w) = \sum_n [y_n - (w_0 + w^T x_n)]^2$$

If we write w_0 equation found above:

$$RSS(w) = \sum_n [y_n - \bar{y} + w^T \bar{x} - w^T x_n]$$

Where $y_n - \bar{y} = y_c$ and $x_n - \bar{x} = x_c$:

$$\begin{aligned}
RSS(w) &= \sum_n [y_c - w^T x_c]^2 = \sum_n (y_c - w^T x_c)(y_n - x_c^T w) \\
&= \sum_n w^T x_c x_c^T w - 2y_c x_c^T w + const = w^T (\sum_n x_c x_c^T) w - 2(\sum_n y_n x_c^T) w + const \\
RSS(w) &= [w^T X_c^T X_c w - 2(X_c^T Y_c)^T w] \\
\frac{\partial RSS(w)}{\partial w} &= X_c^T X_c w - X_c^T Y_c = 0 \\
w^{LMS} &= (X_c^T X_c)^{-1} X_c^T Y_c
\end{aligned}$$

1.2.2 Conditional Gaussian

$[X, Y]$ is multivariate Gaussian distributed.

1.3 Kernel Density Estimation

Random variables X_1, \dots, X_n are i.i.d. sampled from $f(x)$ function.

$$E_{X_1, \dots, X_n}[\hat{f}(x)] = E\left[\frac{1}{nh} \sum_{i=1}^n K\left(\frac{x-t}{h}\right)\right] = \frac{1}{nh} \sum_{i=1}^n E\left[K\left(\frac{x-t}{h}\right)\right] = \frac{1}{nh} n E\left[K\left(\frac{x-t}{h}\right)\right]$$

$$E_{X_1, \dots, X_n}[\hat{f}(x)] = \frac{1}{h} \int K\left(\frac{x-t}{h}\right) f(t) dt$$

$$f(x-hz) = f(x) + f'(x)(x-hz-x) + \frac{f''(x)}{2!}(x-hz-x)^2 + \dots + \frac{f^k(x)}{k!}(x-hz-x)^k + h_k(x-hz)(x-hz-x)^k$$

$$z = \frac{x-t}{h} \text{ which means } t = x - hz$$

$$f(t) = f(x) + f'(x)(t-x) + \frac{f''(x)}{2!}(t-x)^2 + \dots + \frac{f^k(x)}{k!}(t-x)^k + h_k(t)(t-x)^k$$

$$\lim_{t \rightarrow x} h_k(t) = 0$$

$$E_{X_1, \dots, X_n}[\hat{f}(x)] - f(x) = \frac{1}{h} \int K\left(\frac{x-t}{h}\right) f(t) dt - f(x)$$

$$\begin{aligned}
&= \frac{1}{h} \int K\left(\frac{x-t}{h}\right) \left(f(x) + f'(x)(t-x) + \frac{f''(x)}{2!}(t-x)^2 + \dots + \frac{f^{(k)}(x)}{k!}(t-x)^k + h_k(t)(t-x)^k\right) dt - f(x) \\
&= \frac{1}{h} \left(f(x) \int K\left(\frac{x-t}{h}\right) dt + f'(x) \int (t-x) K\left(\frac{x-t}{h}\right) dt + \frac{f''(x)}{2!} \int (t-x)^2 K\left(\frac{x-t}{h}\right) dt + \dots \right. \\
&\quad \left. + \frac{f^{(k)}(x)}{k!} \int (t-x)^k K\left(\frac{x-t}{h}\right) dt + \int h_k(t)(t-x)^k dt \right) - f(x)
\end{aligned}$$

1.4 MLE for Density Estimation

$$f(x) = \frac{1}{n} \sum_{i=1}^n \frac{1}{h} K\left(\frac{x-X_i}{h}\right)$$

$$\begin{aligned}
P(X_1, \dots, X_n | h) &= \prod_{i=1}^n f(X_i) = \left(\frac{1}{h}\right)^n \left(\sum_{i=1}^n K\left(\frac{x-X_i}{h}\right)\right)^n \\
l(h) &= n \log\left(\frac{1}{h}\right) + n \log\left(\sum_{i=1}^n K\left(\frac{x-X_i}{h}\right)\right)
\end{aligned}$$

If we maximize $l(h)$, h value will be 0, which is degenerate.

2 Nearest Neighbor

2.1 Coordinates of Students

2.1.1 Normalizing the Data

There is two dimensions, which requires us to normalize the data points in two different dimension.

$$\hat{x} = \frac{1}{n} \sum_{i=1}^n x_i = \frac{1}{10} \sum_{i=1}^{10} x_i = 14.1$$

$$\hat{y} = \frac{1}{n} \sum_{i=1}^n y_i = \frac{1}{10} \sum_{i=1}^{10} y_i = 21.1$$

Standard deviations:

$$\sigma(x) = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (x_i - \hat{x})^2} = 27.27$$

$$\sigma(y) = \sqrt{\frac{1}{n-1} \sum_{i=1}^n (y_i - \hat{y})^2} = 20.12$$

Normalized data:

$$x_i = \frac{x_i - \hat{x}}{\sigma(x)}$$

$$y_i = \frac{y_i - \hat{y}}{\sigma(y)}$$

Mathematics: (0.045,1.0123),(-1.048,0.620),(-0.899,0.950)

Electrical Engineering: (0.740,0.106),(0.889,0.546),(1.138,0.803)

Computer Science: (0.194,-0.444),(1.287,-1.801),(-1.098,-1.740),(-1.247,-0.334)

2.1.2 Prediction

The student is at(9,18), which normalize to (-0.253,-0.114). For K=1 and using L_2 distance metric. L_2 distance:

$$d_2(a, b) = ||a - b|| = ||a - b||_2 = \sqrt{\sum_{i=1}^d (a_i - b_i)^2}$$

In this case, the distances from each student (3 mathematic, 3 EE and 4 CS respectively):

1.175, 1.08, 1.24, 1.01, 1.31, 1.66, 0.55, 2.28, 1.59, 1.01

The nearest neighbor is CS student(0.55). I predict the new student is CS major.

When K=3, we choose 0.55 (CS), 1.01(CS), and 1.01(EF).Majority is CS. I predict new student is CS major.

For K=1 and using L_1 distance metric. L_1 distance:

$$d_1(a, b) = ||a - b||_1 = \sum_{i=1}^d |a_i - b_i|$$

In this case, the distances with same sequence:

1.43, 1.52, 1.70, 1.21, 1.80, 2.30, 0.77, 3.22, 2.20, 1.21

The nearest student is CS student (0.77). Prediction is the new student is CS major.

When K=3, we choose 0.77(CS), 1.21(CS) , and 1.21(EE). Majority vote is CS. Prediction is new student is CS major.

The results of 4 different prediction methods are same.

2.2 D-Dimensional KNN

2.2.1 $p(x)$

In space, there are N data points, where $\sum_c N_c = N$. In sphere, there are K data points, where $\sum_c K_c = J$. The volume is V.

$$p(x|Y = c) = \frac{K_c}{N_c V}$$

$$p(Y = c) = \frac{N_c}{N}$$

$$p(x) = \sum_c p(x|Y = c)p(Y = c) = \sum_c \frac{K_c}{N_c V} \frac{N_c}{N} = \sum_c \frac{K_c}{V N} = \frac{K}{V N}$$

Bayes rule:

$$p(Y = c|x) = \frac{p(x|Y = c)p(Y = c)}{p(x)} = \frac{\frac{K_c}{V N}}{\frac{K}{V N}} = \frac{K_c}{K}$$

3 Additive or Dropout Noising as Regularization

4 Decision Tree

4.1 First Level Selection

We want to maximize 'Gain' $H[Y] - H[Y|X]$, where $H[Y]$ is fixed. So, we will minimize $H[Y|X]$. Let Y denote yes and N no.

$$\begin{aligned} H[Y|X] &= \sum_k P(X = a_k) H(Y|X = a_k) \\ &= - \sum_k P(X = a_k) [p(Y|X) \log(P(Y|X)) + P(N|X) \log(P(N|X))] \end{aligned}$$

Let's see the gain of "Outlook" feature.

$$\begin{aligned} H[Y|X] = & -P(sunny)[P(Y|sunny)\log(P(Y|sunny))+P(N|sunny)\log(P(N|sunny))] \\ & -P(overcast)[P(Y|overcast)\log(P(Y|overcast))+P(N|overcast)\log(P(N|overcast))] \\ & -P(rain)[P(Y|rain)\log(P(Y|rain))+P(N|rain)\log(P(N|rain))] = 0.48017// \end{aligned}$$

$$\begin{aligned} \text{Temperature feature:} // H[Y|X] = & -P(hot)[P(Y|hot)\log(P(Y|hot))+P(N|hot)\log(P(N|hot))] \\ & -P(mild)[P(Y|mild)\log(P(Y|mild))+P(N|mild)\log(P(N|mild))] \\ & -P(cool)[P(Y|cool)\log(P(Y|cool))+P(N|cool)\log(P(N|cool))] = 0.631501 \end{aligned}$$

$$\begin{aligned} \text{Humidity feature:} // H[Y|X] = & -P(high)[P(Y|high)\log(P(Y|high))+P(N|high)\log(P(N|high))] \\ & -P(normal)[P(Y|normal)\log(P(Y|normal))+P(N|normal)\log(P(N|normal))] = \\ & 0.54651 \end{aligned}$$

$$\begin{aligned} \text{Wind feature:} // H[Y|X] = & -P(strong)[P(Y|strong)\log(P(Y|strong))+P(N|strong)\log(P(N|strong))] \\ & -P(weak)[P(Y|weak)\log(P(Y|weak))+P(N|weak)\log(P(N|weak))] = 0.618397 \end{aligned}$$

Outlook feature minimized $H[Y|X]$ value. First level selector is Outlook value.

4.1.1 Second Level Selection

After first level selection of the Decision Tree, there are 3 subparts. We don't need to split 'Overcast' since all data belongs to Y variable.

Section Sunny: we will try all features. Temperature:

$$\begin{aligned} H[Y|X] = & -P(hot)[P(Y|hot)\log(P(Y|hot))+P(N|hot)\log(P(N|hot))] \\ & -P(mild)[P(Y|mild)\log(P(Y|mild))+P(N|mild)\log(P(N|mild))] \\ & -P(cool)[P(Y|cool)\log(P(Y|cool))+P(N|cool)\log(P(N|cool))] \\ = & -0.4(0\log(0)+1\log(1))-0.4(0.5\log(0.5)+0.5(\log(0.5)))-.2(1\log(1)+0\log(0)) = 0.277258872 \end{aligned}$$

Humidity:

$$\begin{aligned} H[Y|X] = & -P(high)[P(Y|high)\log(P(Y|high))+P(N|high)\log(P(N|high))] \\ & -P(normal)[P(Y|normal)\log(P(Y|normal))+P(N|normal)\log(P(N|normal))] \end{aligned}$$

$$= -0.6(0\log(0) + 1\log(1)) - 0.4(1\log(1) + 0(\log(0))) = 0$$

Wind:

$$\begin{aligned} H[Y|X] &= -P(strong)[P(Y|strong)\log(P(Y|strong))+P(N|strong)\log(P(N|strong))] \\ &\quad -P(weak)[P(Y|weak)\log(P(Y|weak)) + P(N|weak)\log(P(N|weak))] \\ &= -0.4(0.5\log(0.5)+0.5\log(0.5))-0.6(0.33\log(0.33)+0.67(\log(0.67))) = 0.659167373 \end{aligned}$$

Humidity gives the minimum entropy, which will result in maximum Gain. For section Sunny, second selector is Humidity.

Section 'Rain'

Temperature:

$$\begin{aligned} H[Y|X] &= -P(hot)[P(Y|hot)\log(P(Y|hot)) + P(N|hot)\log(P(N|hot))] \\ &\quad -P(mild)[P(Y|mild)\log(P(Y|mild)) + P(N|mild)\log(P(N|mild))] \\ &\quad -P(cool)[P(Y|cool)\log(P(Y|cool)) + P(N|cool)\log(P(N|cool))] \\ &= -0[P(Y|hot)\log(P(Y|hot))+P(N|hot)\log(P(N|hot))]-0.6(0.67\log(0.67)+0.33(\log(0.33)))-.4(0.5\log \end{aligned}$$

Humidity:

$$\begin{aligned} H[Y|X] &= -P(high)[P(Y|high)\log(P(Y|high))+P(N|high)\log(P(N|high))] \\ &\quad -P(normal)[P(Y|normal)\log(P(Y|normal))+P(N|normal)\log(P(N|normal))] \\ &= -0.4(0.5\log(0.5)+0.5\log(0.5))-0.6(0.5\log(0.5)+0.5(\log(0.5))) = 0.659167373 \end{aligned}$$

Wind:

$$\begin{aligned} H[Y|X] &= -P(strong)[P(Y|strong)\log(P(Y|strong))+P(N|strong)\log(P(N|strong))] \\ &\quad -P(weak)[P(Y|weak)\log(P(Y|weak)) + P(N|weak)\log(P(N|weak))] \\ &= -0.4(0\log(0) + 1\log(1)) - 0.6(1\log(1) + 0(\log(0))) = 0 \end{aligned}$$

For section 'Rain', 'Wind' variable gives the minimum entropy. Second selector should be 'Wind'.

4.1.2 Gini Index vs Cross-Entropy

Gini Index: $\sum_{k=1}^K p_k(1 - p_k)$ and Cross-Entropy: $-\sum_{k=1}^K p_k \log_2 p_k$

$$\sum_{k=1}^K p_k(1 - p_k) \leq -\sum_{k=1}^K p_k \log_2 p_k$$

$$\sum_{k=1}^K [1 - p_k + \log_2 p_k] \leq 0$$

If all possible components are less than 0 (or equal), sum will be less than 0. $0 \leq p_k \leq 1$

$$\frac{\partial[1 - p_k + \log_2 p_k]}{\partial p_k} = -1 + \frac{1}{p_k \log 2} = -1 + \frac{1}{0.693 p_k}$$

Considering p_k will be between 0 and 1, this function monotonously increasing between 0 and 1. So, the maximum will be given at 1.

When $p_k = 1$

$$[1 - p_k + \log_2 p_k] = 0$$

For any value of p_k less than 1, this function will give a negative value. This shows sum of $[1 - p_k + \log_2 p_k]$ function of any possible discrete probability functions can't be greater than 0.

5 Programming

5.1 Outliers in Linear Regression

2016/CSCI 567/HW1/yeni/HW1/Outlier.jpg

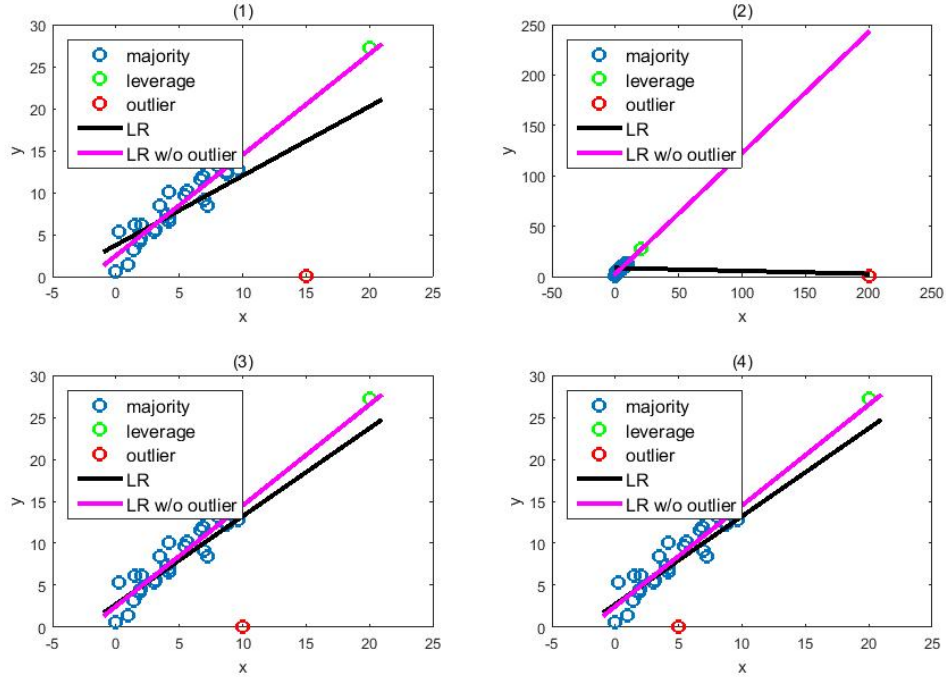


Figure 1: Outliers

5.1.1 Weight Decay

Weight decay on parameters can reduce the effect of the outlier sample in some cases. We can see that, in Figure 1, 3, and 4, fitted lines has less effect of the outlier sample when λ is larger. But in Figure 2, even though λ gets larger, outlier sample's effect can be reduced significantly. I think weight decay on parameters can't be used to reduce the effect of outliers since they reduce the effect of noise but we can't put noises and outliers in the same category.

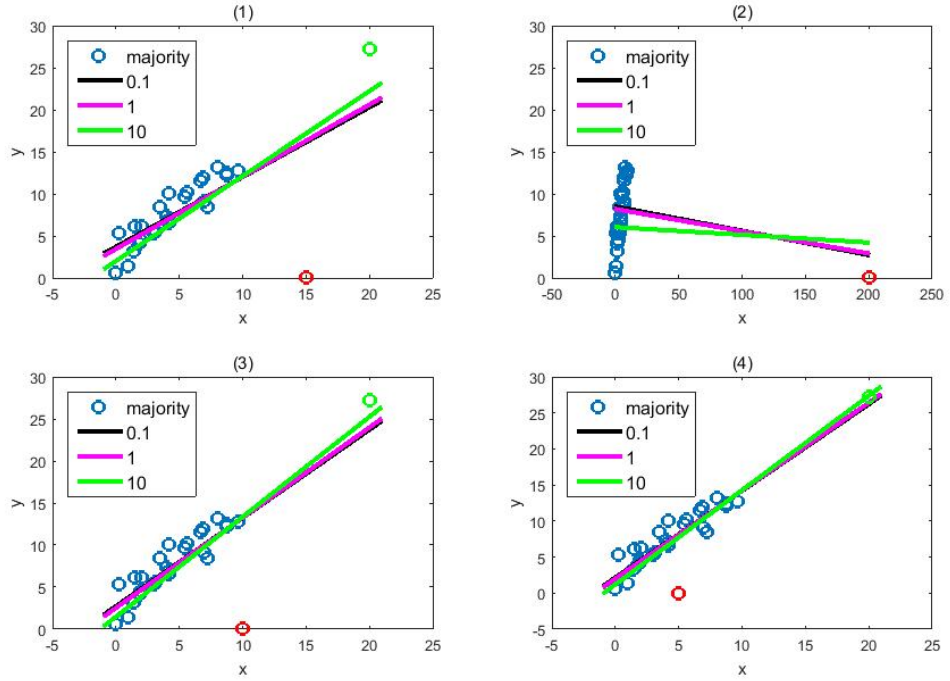


Figure 2: Ridge Regression

5.1.2 Laplace distribution

In linear regression model with Gaussian distribution, the points away from centroid of the data points has a very large effect. Laplace distribution may reduce this effect with correct parameters.

5.1.3 Cook's Distance- Studentized - Leverage Score

Outlier sample can't be chosen using any of these metrics since in Figure 3 and 4 they all select the leverage point instead of the outlier. Cook's Distance chooses the outlier in Figure 1 and 2.

2016/CSCI 567/HW1/yeni/HW1/Cook.jpg

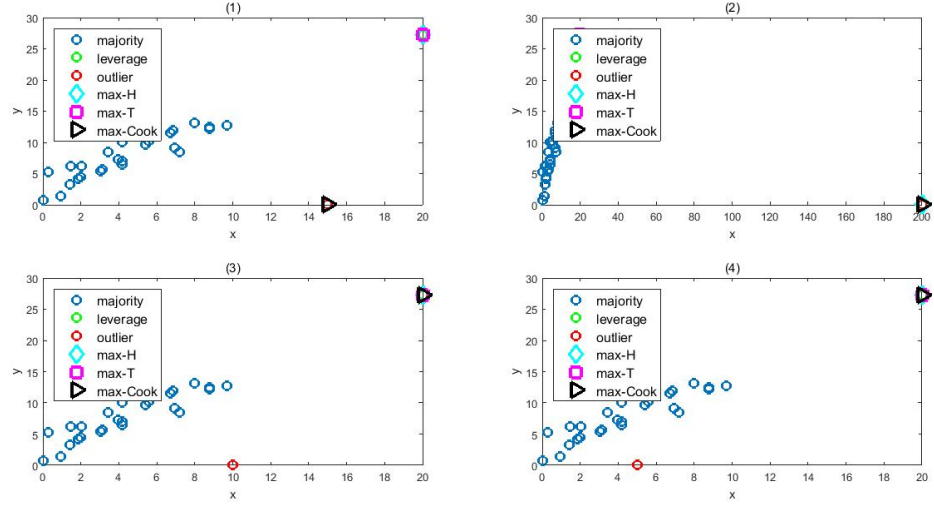


Figure 3: Selecting influential samples

5.2 Classification

5.2.1 kNN Classifier Results

k	Training Acc	Validation Acc	Test Acc
1	0.5835	0.6170	0.5674
3	0.9408	0.6181	0.6084
5	0.9173	0.6287	0.5442
7	0.9511	0.6404	0.5795
9	0.9629	0.6330	0.5879
11	0.9598	0.6245	0.5619
13	0.9614	0.6181	0.5647
15	0.9588	0.6181	0.5498

5.2.2 Decision Tree and Performance Comparison

min Leaf	Split Criteria	Training Acc	Validation Acc	Test Acc
1	Cross Entropy	0.9829	0.6617	0.6567
2	Cross Entropy	0.9829	0.6617	0.6567
3	Cross Entropy	0.9829	0.6617	0.6567
4	Cross Entropy	0.9829	0.6617	0.6567
5	Cross Entropy	0.9787	0.6532	0.6567
6	Cross Entropy	0.9781	0.6404	0.6567
7	Cross Entropy	0.9781	0.6404	0.6567
8	Cross Entropy	0.9781	0.6404	0.6567
9	Cross Entropy	0.9685	0.6883	0.6149
10	Cross Entropy	0.9685	0.6883	0.6149
1	Gini Index	0.9793	0.6489	0.6474
2	Gini Index	0.9793	0.6489	0.6474
3	Gini Index	0.9793	0.6489	0.6474
4	Gini Index	0.9799	0.7043	0.6474
5	Gini Index	0.9787	0.6957	0.6474
6	Gini Index	0.9781	0.6830	0.6474
7	Gini Index	0.9781	0.6404	0.6567
8	Gini Index	0.9736	0.6404	0.6567
9	Gini Index	0.9594	0.6883	0.6149
10	Gini Index	0.9594	0.6883	0.6149

In the provided dataset, Decision Tree is giving a much better accuracy on Training data. But on validation and testing data, accuracies are close to each other even though Decision Tree is giving a better result in general.

5.2.3 Decision Boundary

2016/CSCI 567/HW1/yeni/HW1/k1.jpg




Figure 4: $K = 1$

2016/CSCI 567/HW1/yeni/HW1/k5.jpg




Figure 5: $K = 5$

2016/CSCI 567/HW1/yeni/HW1/k15.jpg

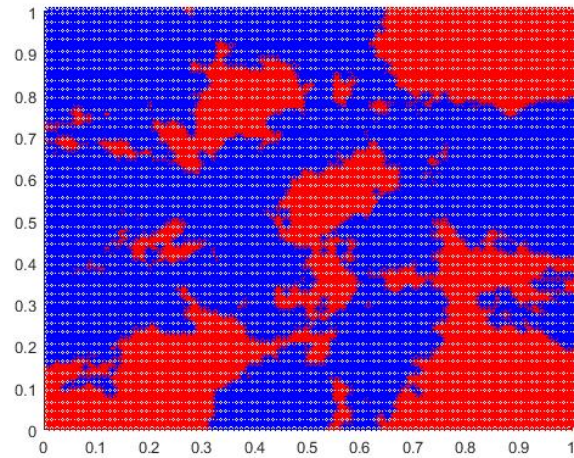


Figure 6: $K = 15$

2016/CSCI 567/HW1/yeni/HW1/k25.jpg




Figure 7: $K = 25$

While k increases, the decision boundary takes a more flexible and smooth shape. When $k = 1$ decision boundaries only depend on one data point, so it makes a discontinuous decision boundary, with sharp points.

On the other hand, $k = 25$ has very large areas that belong to the same class, which is way smoother.