

CSCI 567-Machine Learning Assignment-2

1 Naive Bayes

1.1 Priors and Conditionals

Priors of each dependent variable is:

$$P(no) = \sum_{password} P(no|password)P(password)$$

$$P(no) = P(password = 240)P(no|password = 240) + P(password = 343)P(no|password = 343)$$
$$P(no) = \frac{8}{9}$$

$$P(yes) = \sum_{password} P(yes|password)P(password) = \frac{1}{9}$$

Let's define each x_i : Number i is in password

$$\begin{array}{ll} P(x_0|y = no) = 0.5 & P(x_0|y = yes) = 0.5 \\ P(x_1|y = no) = 0.01 & P(x_1|y = yes) = 0.5 \\ P(x_2|y = no) = 0.5 & P(x_2|y = yes) = 0.5 \\ P(x_3|y = no) = 0.5 & P(x_3|y = yes) = 0.5 \\ P(x_4|y = no) = 1 & P(x_4|y = yes) = 0.5 \end{array}$$

1.2 Password 031

Let N_{x_i} be the number of i 's in password.

$$P(y|x) = \frac{P(x|y)P(y)}{P(x)} = \prod P(x|y)^{N_x} P(y)$$

Two different class:

$$P(yes|031) = \prod P(x_i|yes)^{N_{x_i}} P(yes) = P(x_0|yes)P(x_3|yes)P(x_1|yes)P(yes)$$

$$P(yes|031) = 0.5 * 0.5 * 0.5 * \frac{1}{9} = 0.0134$$

$$P(no|031) = \prod P(x_i|no)^{N_{x_i}} P(no) = P(x_0|no)P(x_3|no)P(x_1|no)P(no)$$

$$P(no|031) = 0.5 * 0.5 * 0.01 * \frac{8}{9} = 0.0022$$

Since $P(yes|031)$ is larger than $P(no|031)$, Naive Bayes Classifier will classify 031 password as "Yes".

1.3 Logistic Regression is Discriminative Version of Naive Bayes

Logistic Regression and Naive Bayes are both classifiers. Naive Bayes Classifier aims to model the joint probability $P(x, y)$ and thus maximizes the joint likelihood $\sum \log P(x_n, y_n)$.

Both Logistic Regression and NBC are linear classifiers that decides on the decision boundary between classes. Naive Bayes' decision boundary is set by $w_0 + \sum_k z_k w_k$ and Logistic Regression's decision boundary is $w_0 + \sum_k w_k x_k$.

To get these decision boundaries, Naive Bayes and other Generative models, model $P(x|y)$ and $P(y)$ and Naive Bayes uses the Bayes' rule to get the joint probability.

$$f(x) = \arg \max_y P(x|y)P(y)$$

On the other hand, Logistic Regression requires only specifying a model for the conditional distribution $P(y|x)$ and maximizes this conditional likelihood $\sum_n \log P(y_n|x_n)$.

Logistic regression sets all the weights together such that the linear decision function tends to be high for positive classes and low for negative.

$$f(x) = \arg \max_y P(y|x)$$

1.4 Assumption of Naive Bayes Classifier

Naive Bayes Classifier sets each features weight independently, based on how much it correlates with the label. The Naive Bayes classifier does this by making a conditional independence assumption that dramatically reduces the number of parameters to be estimated when modeling $P(X|Y)$.

Because this conditional independence assumption is violated when there is correlated features, Logistic Regression outperforms Naive Bayes Classifier in most cases, which makes it more generic than NBC.

2 Generative Model and Discriminative Model

2.1 MLE of Parameters

$$P(D_i|\Pi, \mu_0, \mu_1, S_0, S_1) = [(1 - \Pi) \frac{1}{\sqrt{2\pi}S_0} \exp(-\frac{1}{2S_0^2}(x_i - \mu_0)^2)]^{t_{n0}} \\ ((\Pi) \frac{1}{\sqrt{2\pi}S_1} \exp(-\frac{1}{2S_1^2}(x_i - \mu_1)^2))^{t_{n1}}$$

$$\log P(D|\Pi, \mu_0, \mu_1, S_0, S_1) = \sum_{n=1}^N t_{n0} [\log(1-\Pi) - \frac{1}{2} \log(2\pi) - \log S_0 - \frac{1}{2S_0^2} (x_n - \mu_0)^2] + \\ \sum_{n=1}^N t_{n1} [\log \Pi - \frac{1}{2} \log(2\pi) - \log S_1 - \frac{1}{2S_1^2} (x_n - \mu_1)^2]$$

Derivative for Π :

$$\frac{\partial \log P(D|\Pi, \mu_0, \mu_1, S_0, S_1)}{\partial \Pi} = N_1 \frac{1}{\Pi} - N_0 \frac{1}{1 - \Pi} = 0$$

$$\Pi = \frac{N_1}{N_0 + N_1}$$

Derivative for μ_0 and μ_1 :

$$\frac{\partial \log P(D|\Pi, \mu_0, \mu_1, S_0, S_1)}{\partial \mu_0} = \frac{\sum t_{n0} x_n}{S_0} - \frac{N_0 \mu_0}{S_0} = 0$$

$$\mu_0 = \frac{\sum t_{n0} x_i}{N_0}$$

$$\mu_1 = \frac{\sum t_{n1} x_i}{N_1}$$

Derivative for S_0 and S_1 :

$$\frac{\partial \log P(D|\Pi, \mu_0, \mu_1, S_0, S_1)}{\partial S_0} = -\frac{N_0}{S_0} + \frac{1}{2S_0^2} \sum t_{n0} (x_n - \mu_0)^2 = 0$$

$$S_0 = \frac{\sum t_{n0} (x_n - \mu_0)^2}{N_0}$$

$$S_1 = \frac{\sum t_{n1} (x_n - \mu_1)^2}{N_1}$$

2.2 Posterior Probability

$$p(y = 1|x, \pi, \mu_0, \mu_1, S_0, S_1) = \frac{p(y = 1|\pi)p(x|y, \mu_1, S_1)}{p(y = 1|\pi)p(x|\mu_1, S_1) + p(y = 0|\pi)p(x|\mu_0, S_0)}$$

This has the form of:

$$P(y = 1|x) = \frac{P(y = 1)p(x|y = 1)}{P(y = 1)p(x|y = 1) + p(y = 0)p(x|y = 0)}$$

Divide both nominator and denominator by nominator:

$$\begin{aligned} &= \frac{1}{1 + \frac{p(y=0)p(x|y=0)}{P(y=1)p(x|y=1)}} = \frac{1}{1 + \exp(\log \frac{p(y=0)p(x|y=0)}{P(y=1)p(x|y=1)})} \\ &= \frac{1}{1 + \exp(\log \frac{p(y=0)}{p(y=1)} + \sum_n \log \frac{p(x_n|y=0)}{p(x_n|y=1)})} \end{aligned}$$

Given the parameters in the question and $S_0 = S_1 = S$:

$$\sum_n \log \frac{p(x_n|y = 0)}{p(x_n|y = 1)} = \sum_n \log \frac{\frac{1}{\sqrt{2\pi}S_n} \exp(-\frac{(x_n - \mu_{n0})^2}{2S_n^2})}{\frac{1}{\sqrt{2\pi}S_n} \exp(-\frac{(x_n - \mu_{n1})^2}{2S_n^2})} = \sum_n (\frac{\mu_{n0} - \mu_{n1}}{S_n^2} x_n + \frac{\mu_{n1}^2 - \mu_{n0}^2}{2S_n^2})$$

Put in equation:

$$P(y = 1|x) = \frac{1}{1 + \exp(\log \frac{1-\pi}{\pi} + \sum_n (\frac{\mu_{n0} - \mu_{n1}}{S_n^2} x_n + \frac{\mu_{n1}^2 - \mu_{n0}^2}{2S_n^2}))}$$

In this case:

$$\begin{aligned} w_0 &= \log \frac{1-\pi}{\pi} + \sum_n \frac{\mu_{n1}^2 - \mu_{n0}^2}{2S_i^2} \\ w_i &= \frac{\mu_{i0} - \mu_{i1}}{S_i^2} \end{aligned}$$

2.3 Poisson Mixture Distribution

Starting from this:

$$\begin{aligned} P(y = 1|x) &= \frac{1}{1 + \exp(\log \frac{p(y=0)}{p(y=1)} + \sum_n \log \frac{p(x_n|y=0)}{p(x_n|y=1)})} = \frac{1}{1 + \exp(\log \frac{1-\pi}{\pi} + \sum_i \log \frac{\exp(-\lambda_{0i})\lambda_{0i}^{x_i}}{\exp(-\lambda_{1i})\lambda_{1i}^{x_i}})} \\ &= \frac{1}{1 + \exp(\log \frac{1-\pi}{\pi} + \sum_i (\lambda_{1i} - \lambda_{0i})x_i (\log \lambda_{0i} - \log \lambda_{1i}))} \\ w_0 &= \log \frac{1-\pi}{\pi} \\ w_i &= (\lambda_{1i} - \lambda_{0i})(\log \lambda_{0i} - \log \lambda_{1i}) \end{aligned}$$

3 Multinomial Logistic Regression

4 Programming Questions

Method	Data Type	Training Acc	Test Acc
Batch Gradient	raw	0.8692	0.8817
Batch Gradient	standardized	0.9300	0.9261
Newton's Method	raw	0.8935	0.8922
Newton's Method	standardized	0.8935	0.8922
glmfit	raw	0.9357	0.9265
glmfit	standardized	0.9357	0.9265

4.0.1 Chosen 20 Features

The feature IDs chosen according to Mutual Information values:

[52, 53, 56, 21, 55, 7, 16, 57, 25, 19, 24, 5, 27, 17, 3, 26, 23, 6, 11, 2]

Using these 20 features on glmfit function: Training accuracy = 0.8914
and Testing Accuracy = 0.8926

4.1 Generative Model and Discriminative Model

4.1.1 MLE Parameters When Variances Are Independent

Component Proportions = ([0.2941,0.2679,0.4379])

μ Parameters = ([1.538,0.0354; -1.593,1.596; -1.486,-1.408])

Σ Parameters =

First Component	Second Component	Third Component
0.9988 0.0102	1.0197 -0.4525	1.0280 0.5110
0.0102 3.9734	-0.4525 1.0105	0.5110 1.1234

Testing Accuracy = 0.9007

4.1.2 MLE Parameters When Variances Are Equal

Component Proportions = ([0.293,0.275,0.430])

μ Parameters = ([1.549,0.0435;-1.759,-0.739;-1.385,0.0282])

Σ Parameters =

0.977 -0.004

-0.004 3.331

Testing Accuracy = 0.7600

4.1.3 Multinomial Logistic Model

Testing Accuracy = 0.8846

4.1.4 Plot and Analysis

We can see that the the best decision boundaries are given by Gaussian Mixture Model with different variance. And the worst one is given by Gaussian Mixture Model with same variance.

It is obvious that variances of three Gaussian models are different from each other. Even though green and blue colored data points have a similar variance, red colored data points' distribution has much larger variance. Since we made the same variance assumption in second model, it is reasonable to get not a very good decision boundary.

On the other hand, logistic regression has made a very good classification and draw a good decision boundary between classes even though it couldn't use the advantage of being nonlinear as different variance Gaussian Mixture Model did.

4.1.5 Subset Training Data

3 plots below shows the testing accuracy w.r.t. training data sizes for three models.

4.2 Practical Logistic Regression on Toy Data

4.2.1 Discretization

The best model is chosen according to the best cross-validation accuracy. Leave one out technique is used.

Data	Best Discretization	Training Acc	Test Acc	Heldout Acc
1	16^2	0.9950	0.9737	0.9775
2	16^2	0.9975	0.9949	0.9925
3	16^2	1	0.9950	1
4	4^2	0.9800	0.9539	0.9675

4.2.2 l_2 norm regularization)

For Data1:

Type	$\lambda = 1$	$\lambda = 0.1$	$\lambda = 0.01$	$\lambda = 0.001$
Training Accuracy	0.9700	0.9750	0.9775	0.9775
Heldout Accuracy	0.9700	0.9700	0.9700	0.9700
Testing Accuracy	0.9559	0.9583	0.9611	0.9612

For Data2:

Type	$\lambda = 1$	$\lambda = 0.1$	$\lambda = 0.01$	$\lambda = 0.001$
Training Accuracy	0.9850	0.9850	0.9825	0.9825
Heldout Accuracy	0.9825	0.9825	0.9800	0.9800
Testing Accuracy	0.9833	0.9871	0.9915	0.9915

For Data3:

Type	$\lambda = 1$	$\lambda = 0.1$	$\lambda = 0.01$	$\lambda = 0.001$
Training Accuracy	0.9800	0.9850	0.9900	0.9950
Heldout Accuracy	0.9800	0.9800	0.9800	0.9800
Testing Accuracy	0.9741	0.9760	0.9779	0.9783

For Data4:

Type	$\lambda = 1$	$\lambda = 0.1$	$\lambda = 0.01$	$\lambda = 0.001$
Training Accuracy	0.9525	0.9800	0.9800	0.9825
Heldout Accuracy	0.9425	0.9675	0.9550	0.9625
Testing Accuracy	0.9425	0.9539	0.9520	0.9587

4.2.3 Visualization

For regularized terms; for Data2 λ chosen as 1 and for Data3 λ is chosen as 0.1.// For unregularized models for data2 and data3 are both 16^2 discretizations.